

## Exercise 9

**Task 9.1** (Stencil algorithm on Xeon Phi). In this task you will work with a 9-point stencil algorithm as it is known from calculating partial differential equations. The algorithm runs over a 2d grid and computes a weighted sum of each entry with all its neighbors.

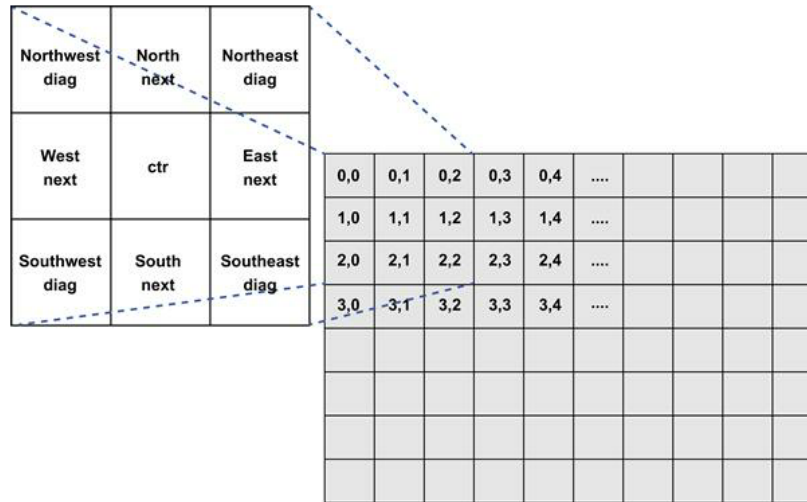


Figure 9.1: 9 point stencil

The results are written in a second array to avoid conflicts. After each iteration input- and output-array are swapped. The C-code of the algorithm can be found in the following.

```

1  void
2  stencil9pt_omp (REAL *finp, REAL *foutp,
3                  int width, int height,
4                  REAL ctr, REAL next, REAL diag, int count)
5  {
6      REAL *fin = finp;
7      REAL *fout = foutp;
8      int i, x, y;
9
10     for (i=0; i<count; i++) {
11
12
13         for (y=1; y < height-1; y++) {
14             // starting center pt (avoid halo)
15             int c = 1 + y*WIDTHHP+1;
16             // offsets from center pt.
17             int n = c-WIDTHHP;
18             int s = c+WIDTHHP;
19             int e = c+1;
20             int w = c-1;
21             int nw = n-1;
22             int ne = n+1;
23             int sw = s-1;
24             int se = s+1;
25
26
27             for (x=1; x < width-1; x++) {
```

```
28         fout[c] = diag * fin[nw] +
29                 diag * fin[ne] +
30                 diag * fin[sw] +
31                 diag * fin[se] +
32                 next * fin[w] +
33                 next * fin[e] +
34                 next * fin[n] +
35                 next * fin[s] +
36                 ctr * fin[c];
37
38         // increment to next location
39         c++;n++;s++;e++;w++;nw++;ne++;sw++;se++;
40     }
41 }
42 REAL *ftmp = fin;
43 fin = fout;
44 fout = ftmp;
45 }
46 return;
47 }
```

The kernel should now be prepared for execution on the Xeon Phi.

- (a) The compiler reports that he assumes a vector dependency when trying to vectorize the only loop candidate. What loop is this candidate? What can you do in order to resolve this issue?
- (b) After the code vectorizes it needs to be parallelized. Where and what appropriate OpenMP pragmas must be placed in order to do worksharing?

**Task 9.2** (Ellpack Storage Format). In this task the Ellpack storage format, known from the lecture slides, is investigated. As test case we are employing a matrix-vector-multiplication. A framework for this task is available in *05task\_crs.zip* on moodle. It is compileable with *g++ 05task\_ellpack.cpp -o ellpack -fopenmp* and can be executed. Of course the integrated verification of this file fails at the moment.

- (a) First the input data has to be converted to the Ellpack format. To that end, complete the function *createELLPACKSystem* within the file *05task\_ellpack.cpp*. Comments within the skeleton of the function indicate where you have to insert your code.
- (b) Now the actual method for the calculation must be implemented. A code skeleton for the function *multiplyELLPACK* is already included in *05task\_crs.cpp*. You need not to implement an optimised or vectorized kernel.